

Learning to Look at the Other Side: A Semantic Probing Study of Word Embeddings in LLMs with Enabled Bidirectional Attention

Zhaoxin Feng and Jianfei Ma and Emmanuele Chersoni
and Xiaojing Zhao and Xiaoyi Bao

Chinese and Bilingual Studies, The Hong Kong Polytechnic University
{zhaoxinbetty.feng,jian-fei.ma,xiaojing.zhao,xiaoyi.bao}@connect.polyu.hk,
emmanuele.chersoni@polyu.edu.hk

Abstract

Autoregressive Large Language Models (LLMs) demonstrate exceptional performance in language understanding and generation. However, their application in text embedding tasks has been relatively slow, along with the analysis of their semantic representation in probing tasks, due to the constraints of the unidirectional attention mechanism.

This paper aims to explore whether such constraints can be overcome by enabling bidirectional attention in LLMs. We tested different variants of the Llama architecture through additional training steps, progressively enabling bidirectional attention and unsupervised/supervised contrastive learning.

Our results show that bidirectional attention improves the LLMs’ ability to represent subsequent context but weakens their utilization of preceding context, while contrastive learning training can help to maintain both abilities¹.

1 Introduction

Decoder-only LLMs using autoregressive pretraining have achieved superior performance across language understanding and generation tasks, causing a major shift from the previous pretraining-then-finetuning paradigm dominated by encoder-only models (Naveed et al., 2023). However, the community has been relatively slow in adopting them for word, sentence, and document embedding tasks because of their apparent limitations as text encoders, which have been speculated to be due to the lack of bidirectional attention (Qorib et al., 2024; BehnamGhader et al., 2024; Springer et al., 2025). As illustrated in Figure 1, decoder-only LLMs can only access preceding contextual information during inference, resulting in word representations that encode information from the previous context, instead of the entire input sequence.

¹Our code and data are released at: <https://github.com/Zhaoxin-Feng/semantic-probing-2025>.

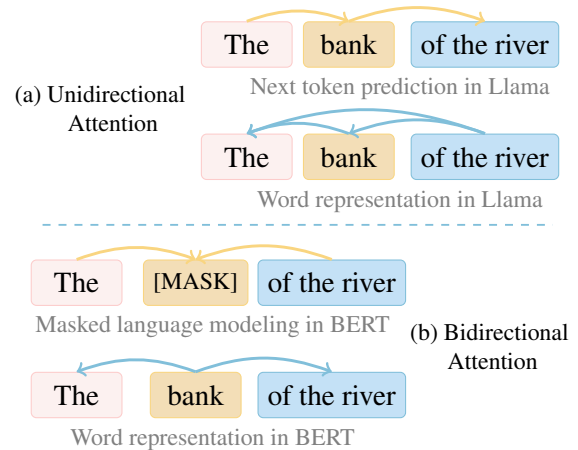


Figure 1: Comparison of attention mechanisms in Llama and BERT models. (a) shows Llama’s unidirectional attention where prediction (orange arrows) and word representation (blue arrows) can only access one side context; (b) shows BERT’s bidirectional attention where masked language modeling allows word representation to access both previous and subsequent context.

This architectural constraint is potentially very limiting in tasks requiring fine-grained modulation of word meanings: while contextualized embeddings from encoder-only models marked significant progress compared to previous generation distributional models (Bommasani et al., 2020; Chronis and Erk, 2020), the availability of the right-hand context might be important to capture subtle meaning nuances, and disambiguate the senses of polysemous words (Zhu et al., 2024; Qorib et al., 2024).

Can this limitation be addressed? After all, the decoder-only architecture enables more efficient learning from all input tokens during pre-training, significantly improving sample efficiency compared to encoder-only counterparts (Clark et al., 2020), and this would be an important advantage if LLMs representations could be adapted to perform better in embedding tasks.

Given the above-mentioned research back-

ground, we propose the following research question: *does bidirectional attention in LLMs enhance the quality of word meaning representations in LLMs? Could they achieve the same quality of embeddings extracted from encoder-only models?*

In this paper, we propose a probing study of different types of LLMs architectures on a pool of *lexical semantic tasks*. Drawing inspiration from recent work on enabling bidirectional attention in autoregressive models (BehnamGhader et al., 2024), we compare the performance of Llama embeddings under the following configurations: i) a base Llama architecture, ii) the architecture in i) after an additional training step to enable bidirectional attention; iii) the architecture in ii), after applying unsupervised/supervised contrastive learning.

Perhaps surprisingly, we found that bidirectional attention in itself does not improve the performance of Llama embeddings on our tasks: while it improves LLMs’ ability to represent the right-hand context of a target word, it also seems to weaken the representation of the left context. Contrastive learning techniques often help the models to maintain both abilities, with Llama architectures getting on par or even outperforming bidirectional, BERT-based baselines on our tasks. Interestingly, we also found that adding bidirectional attention alone exacerbates the *anisotropy* (Ethayarajh, 2019; Cai et al., 2021; Godey et al., 2024) (a condition in which all vectors occupy just a narrow cone in the vector space) in all layers, resulting on average in higher similarity scores between the vectors of randomly sampled words. These findings reveal the potential of decoder-only LLMs in word embedding tasks, offering insights into enhancing LLMs’ representations with bidirectional attention and contrastive learning.

2 Related Work

2.1 Representations of Word Semantics

The representations of word semantics in NLP have undergone a remarkable development in the last two decades. Early methods like distributional semantic models (DSMs) derive semantic representations from statistical patterns of word co-occurrences in large text corpora, assuming that words with similar contexts have similar meanings (Harris, 1954; Schütze, 1992; Bullinaria and Levy, 2012). Later more efficient methods like Word2Vec and GloVe (Mikolov et al., 2013; Pennington et al., 2014) emerged, using neural net-

works to train word embedding representations more compactly, without time-consuming high dimensional sparse data processing and high perplexity algorithms calculation (Pennington et al., 2014; Mikolov et al., 2013).

However, these static vector models struggled with polysemy, as they assigned a single vector to each word regardless of context (Faruqui et al., 2016; Gladkova and Drozd, 2016; Wang et al., 2020b). This limitation was addressed by contextualized embedding models such as ELMo and BERT (Peters et al., 2018; Devlin et al., 2019), in which word vectors are learned as a function of the internal states of the network, such that a word in different sentence contexts determines different activation states and is represented by a distinct vector (Chersoni et al., 2021).

Despite the advantages of representing context-specific meanings, contextualized vectors were shown to have a high level of anisotropy, i.e. they occupy just a narrow cone in the vector space, with the consequence that randomly-sampled words might also get high similarity values (Ethayarajh, 2019) and postprocessing techniques need to be applied to adjust the similarity metrics for anisotropy (Timkey and van Schijndel, 2021).

2.2 Probing Linguistic Features in LLMs

Probing-based methods for analyzing linguistic features have become prevalent for understanding the internal linguistic knowledge of language models (Linzen et al., 2016; Hewitt and Liang, 2019; Liu et al., 2019; Wu et al., 2020; Chersoni et al., 2021; Kauf et al., 2023; Matthews et al., 2024; Wang et al., 2024a; Liu et al., 2024). The core idea involves using a simple diagnostic model (the “probe”, usually a linear classifier) to predict specific linguistic properties (e.g. animacy) from language model’s output representations. If the model succeeds, we can infer that the representations of the language model encode that linguistic knowledge (Chersoni et al., 2021). For instance, Hewitt and Manning (2019) presents a linear classifier to predict a target syntactic structure based on contextualized word representations to measure the syntactic information encoded in language models.

For LLMs, the most prevalent method to probe the linguistic knowledge is to use the representation is the hidden state of the last layer (Neelakantan et al., 2022; BehnamGhader et al., 2024; Springer et al., 2025; Lee et al., 2025). We went for the last layer embedding also for an issue of method-

ological alignment with these approaches. Also, considering that the last hidden layer is processed through all model layers, in theory it contains the richest and most comprehensive context information, and thus our method employs these final layer hidden states to derive contextual representations of target words.

2.3 Autoregressive LLMs as Text Encoders

While ELMo and BERT’s bidirectional attention gives them access to both the left and the right context around the target word, autoregressive LLMs like the GPTs (Radford et al., 2019; Brown et al., 2020) only use the context that comes before the target. Therefore, autoregressive LLMs are often considered sub-optimal for text embedding tasks.

However, recent studies have shown that even under the constraints of causal masking, LLMs are still capable of capturing certain contextual relationships (Muennighoff, 2022; Wang et al., 2024b; BehnamGhader et al., 2024; Springer et al., 2025). Among these studies, both Springer et al. (2025) and BehnamGhader et al. (2024) specifically hypothesize that the limitation of LLMs lies in unidirectional attention: the former proposed “echo embeddings”, a method that feeds a target input sequence twice to a model to allow for the encoding of the context after the target; while the latter proposes an additional training step, called *masked next token prediction*, to enable bidirectional attention in autoregressive models, and introduces further refinements based on unsupervised and supervised contrastive learning techniques, in order to improve the performance in sentence-level tasks. We ground our work in BehnamGhader et al. (2024)’s LLM2Vec framework, using its publicly available models to observe how different training strategies affect the models’ behavior in lexical semantic tasks.

3 Experimental Setup

3.1 Model Selection

Our experiments focused on Llama architectures (Touvron et al., 2023a), and particularly on Sheared-Llama (Xia et al., 2023) and Llama 2 (Touvron et al., 2023b). The former was chosen because it is a structurally-pruned and space-efficient version of the original model (we used the 1.3B version), and it is the closest in terms of parameter size to the most commonly-used bidirectional models (e.g. BERT and RoBERTa). We selected the 7B

version of Llama 2 mainly to check if the trends identified with Sheared-Llama-1.3B were consistently observed in a bigger model.

Besides the base models (i.e. **Sheared-Llama-1.3B** and **Llama2-7B**), for our experiments we tested their variants augmented with the additional training steps of the LLM2Vec framework (BehnamGhader et al., 2024): 1) the **Bi+MNTP** models underwent an additional training via *masked next token prediction*, in which part of the input tokens was masked and the model had to reconstruct them on the basis of left and right context. For the prediction of each masked token, only the logits obtained from the previous token positions were used for computing the loss; 2) the architecture in 1), but with the addition of unsupervised (**SimCSE**) and **Supervised** contrastive learning (Gao et al., 2021) on top of the bidirectional training. BehnamGhader et al. (2024) added this step claiming that an autoregressive LLM with Bi+MNTP could be adequate for word-level tasks (they indeed obtain improved performance on standard benchmarks for POS Tagging and Named Entity Recognition), but contrastive learning is helpful to make their sequence representations a good fit for sentence-level tasks as well. All the augmented models are based on the same Sheared-Llama and Llama 2 architectures, and this gives us the chance to directly compare the effects of each additional training step.

We also used the embeddings from BERT Base and BERT Large (Devlin et al., 2019) as our bidirectional baselines. For more details about our LLMs and the settings of the probe classifiers, the reader can refer to Appendix A.2.

3.2 Extracting Word Representations

We leverage the hidden states from the final layer of LLMs to obtain contextualized representations of the target word for each lexical semantic task. Given a word w within a sequence c , we first extract sequence representations as follows:

$$\mathbf{H} = [\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_n] \in \mathbb{R}^{n \times d} \quad (1)$$

where n denotes the sequence length (number of tokens) and d represents the no. of hidden dimensions (e.g., 2048 in Sheared-Llama-1.3B).

For a word w tokenized into k subwords $\{t_1, t_2, \dots, t_k\}$, let $\mathcal{I} = \{i_1, i_2, \dots, i_k\}$ denote their positional indices in \mathbf{H} . The final word representation \mathbf{v}_w is obtained through:

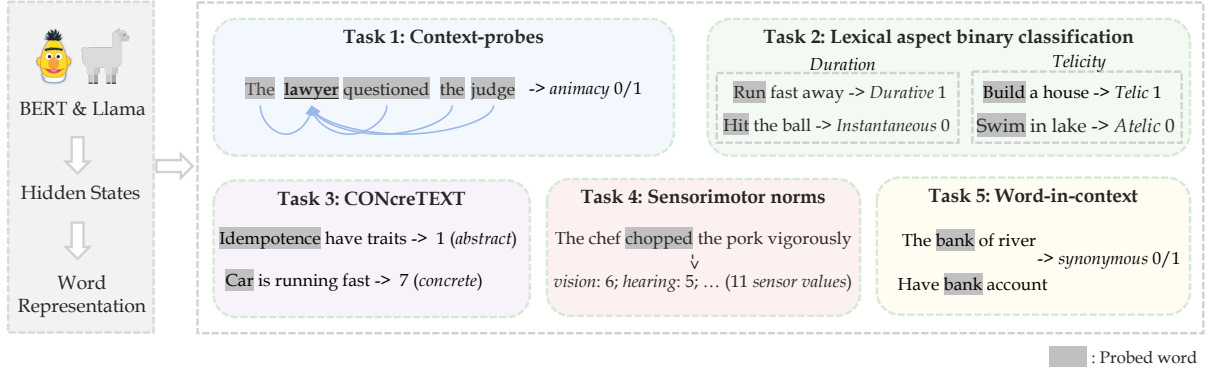


Figure 2: Summary of five semantic probing tasks in our study: Tasks 1, 2, and 5 are classification tasks, with 0 and 1 denoting binary labels; Tasks 3 and 4 are regression tasks, where numbers (eg. 5, 7) indicate continuous values for the target variables. The “probed word” (highlighted) refers to the word whose contextualized representation is extracted for the probing task.

$$\mathbf{v}_w = \frac{1}{|\mathcal{I}|} \sum_{j \in \mathcal{I}} \mathbf{h}_j \quad (2)$$

When $|\mathcal{I}| = 1$ (a single-token word), Equation 2 simplifies to $\mathbf{v}_w = \mathbf{h}_i$. This formulation ensures consistent representation for both single- and multi-token words while preserving contextual information from the final layer.

The first four tasks all involve a single target word in a sentence context, so we directly feed the target word’s representation to a classifier or regression model for prediction. The exception was the **Word-in-Context** (see Section 3.3), which required comparing the meaning of a target word in two different sentences. In this task, we test three types of embeddings-derived features as input to a classifier: their *cosine similarity*, the *absolute values of their element-wise difference*, and their *concatenation*.

3.3 Datasets and Tasks

Our study adopts five probing datasets (Figure 2). First, the **context-probes** dataset by Klafka and Ettinger (2020), which is composed of sets of five-word sentences with a subject-verb-object structure and a binary property annotated for the verb or one of the nouns in the sentence, e.g. animate vs. inanimate for one of the nouns (the subject or the object); dynamic vs. static, and supporting causative-inchoative alternation for verbs. This is an easy task, targeting relatively stable properties of the target words regardless of context, but it can be useful to test the attention patterns of a model: in theory, LLMs with bidirectional attention should be able to solve it *regardless of which token embed-*

ding is fed to the classifier (i.e., even if the target word follows the token embedding, a contextualized embedding from a BERT model should still contain information from the right-hand context), while autoregressive LLMs should struggle more when the input embedding is extracted from a token preceding the target.

A second task, targeting verbs, is **lexical aspect classification** (Metheniti et al., 2022). This dataset contains annotations about telicity and duration for verbs in a sentence context: given a verb embedding, a probing model has to determine whether it is telic/atelic or durative/stative.

Next, we ran two regression tasks using two datasets targeting the semantics of words in context: **CONcreTEXT** (Gregori et al., 2020) and **contextualized sensorimotor norms** (Trott and Bergen, 2022). The former contains annotations about the concreteness of a given word, with mean human ratings on a Likert scale ranging from 1 (totally abstract) to 7 (totally concrete); the latter contains mean human ratings for both verbs and nouns on 11 different sensorimotor domains, such as vision, hearing, etc. For each word, a sensorimotor rating indicates to what extent the corresponding sense is relevant to experiencing the concept referred by the target word in that specific context. In both cases, the goal for a probe model is to predict human ratings using as input the embedding of the target word (Fagarasan et al., 2015; Utsumi, 2018; Thompson and Lupyan, 2018; Li and Summers-Stay, 2019; Chersoni et al., 2020, 2021; Flor, 2024).

For the final task, we employ the **Word-in-Context** (Pilehvar and Camacho-Collados, 2019)

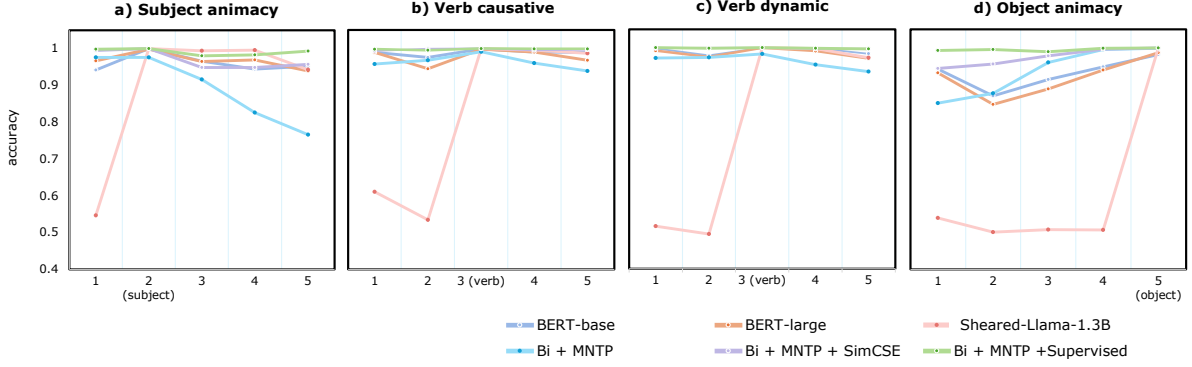


Figure 3: Results of predicting subject animacy, verb causative/dynamic, and object animacy using each word in a sentence as probed words, scores extracted with Sheared-Llama-1.3B and its variants. The horizontal axis represents word indices in sentences (all with identical five-word syntactic structures).

dataset. In each instance, a target word appear in two different sentences, and the probe model has to determine whether the word is being used in the same sense or not. Dataset details are described in Appendix A.1.

We selected the probing tasks to address different types of semantic features for different parts-of-speech (e.g. nouns and verbs), and the tasks demand a different level of contextual sensitivity: while the features of *context-probes* or verb aspect should be stable for a target word across linguistic contexts, the regression tasks and the *Word-in-Context* were explicitly designed to require a deeper understanding of contextual meaning.

As for the probes, we tested both a linear and a non-linear model on top of the embeddings: the former was logistic regression for classification tasks and linear regression for concreteness and norms predictions; the latter was a multilayer perceptron (see Appendix B). In the main text, we only report the results of the Multilayer Perceptron, which achieved higher scores, while the results for the linear models are in Appendix B.2. All tasks were implemented on a single 40GB NVIDIA A100 GPU.

To ensure that our probing methodology actually evaluates the quality of contextual information in embeddings and to assess our selected tasks’ sensitivity to context, we conducted two complementary experiments documented in the Appendix ²: 1) control tasks with randomly sampled labels (Appendix C), and 2) out-of-context evaluations where models processed probed words in isolation rather than complete sentences (Appendix D).

²The experiments were added upon a reviewer’s request.

4 Results and Analysis

The experimental results yield three main findings derived from Tasks 1 and 2, Tasks 3 and 4, and the final task, respectively.

Finding 1: *Bidirectional attention improves the LLMs’ ability to represent subsequent context, but it weakens the utilization of the previous context. Contrastive learning techniques mitigate this trade-off by enhancing the model’s ability to balance contextual understanding in both directions.*

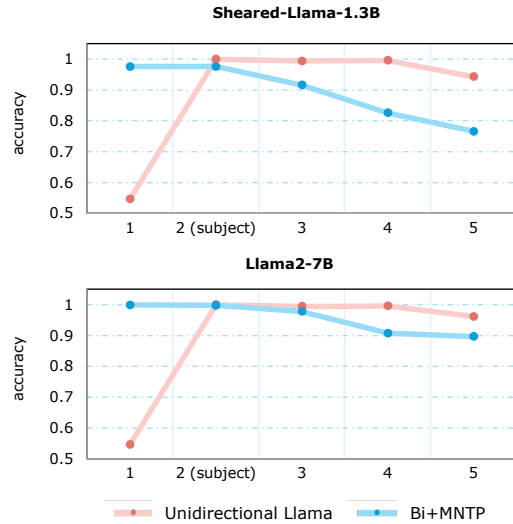


Figure 4: Results of subject animacy subtask in Task 1 by comparing Sheared-Llama-1.3B to Llama2-7B.

Figure 3 shows the results for the *context-probes* dataset, for the base, unidirectional Llama, when the embedding of the probed word comes before the target word, the model has severe difficulty in correctly predicting the properties of the target word. On the other hand, it can be seen that the

performance of the bidirectional baselines is consistently high, regardless of the token embedding. Once bidirectional attention is enabled, however, the score pattern for Llama Bi + MNTP is more aligned with the BERT models.

There is, though, an important difference: by comparing Bi + MNTP and unidirectional Llama in Figure 3 a), we can find that the pink line shows relatively stable accuracy in the latter half, while the blue line exhibits a noticeable decline as the distance between target word and token of the input embedding increases. This indicates that enabling bidirectional attention may also weaken Llama’s inherent ability to “see” the preceding context.

At the same time, if we look at the models enhanced with contrastive learning (purple and green lines), we can notice that the additional training helps Llama to maintain representation quality in all the token positions. The results of the first task thus suggest that the additional training to refine sentence-level representations has positive effects also on the quality of single token embeddings. Figure 4 shows a comparison between Sheared-Llama and Llama 2 in terms of the impact of bidirectional attention, and it can be seen that the larger model is following a similar pattern, although the accuracy scores have a less sharp decrease. Similarly to Sheared-Llama, the application of contrastive learning greatly helps model performance (see Appendix B for more detailed results).

In Figure 5 we can see the scores for the *lexical aspect classification*, where we can observe that, for telicity, the basic versions of Llama are already performing well and on par with the bidirectional baselines, while they fall short of BERT Base in modeling verb duration. As we observed before, bidirectional training alone with Bi + MNTP actually makes the models less accurate, whereas contrastive learning techniques consistently improve their performance. At a glance, it can be noticed that Llama 2 with either contrastive training type is better than both BERT baselines in both subtasks.

Finding 2: *Autoregressive models perform similarly to bidirectional ones on regression probing tasks for norms prediction.*

Moving to the regression tasks, e.g. concreteness prediction on *CONcreTEXT* and modeling of *contextualized sensorimotor norms* (Figure 6), we can immediately see that base Llama models already achieve high correlations with human mean ratings, with no significant disadvantage to the bidirectional competitors. This contradicts previous

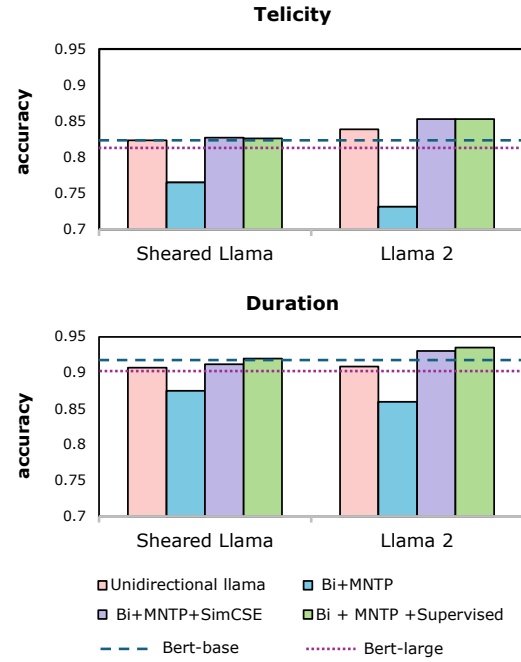


Figure 5: Results of verb telicity and duration (Task 2).

findings claiming that autoregressive LLMs are not optimal when it comes to modeling word semantics (Qorib et al., 2024).

It should not be underestimated, however, the specificity of the task at hand: while most probing tasks presented in the literature focus on discrete distinctions (e.g. grammatical vs. ungrammatical, semantically plausible vs. implausible), there is only limited comparative evidence about LLM performance when it comes to modeling fine-grained, continuous judgements of semantic properties. Interestingly, recent studies reported that LLMs are able to faithfully reproduce human ratings of concreteness and sensory norms via prompting GPT-4 (Xu et al., 2023; Martínez et al., 2025), and thus it is possible that such semantic features are robustly encoded also in other autoregressive architectures (e.g. Llama models).

Once again, Bi + MNTP alone deteriorates embedding quality, and contrastive learning strategies mitigate its negative effects. However, in the concreteness task this does not lead to any consistent improvement over the base models; in the sensorimotor norms task, better correlations can be observed only for the Llama 2 model, and only with the unsupervised contrastive learning strategy.

In addition, it could be noticed about these tasks that in none of the model families size seems to matter too much: Sheared-Llama and BERT Base are often on par or better than the corresponding

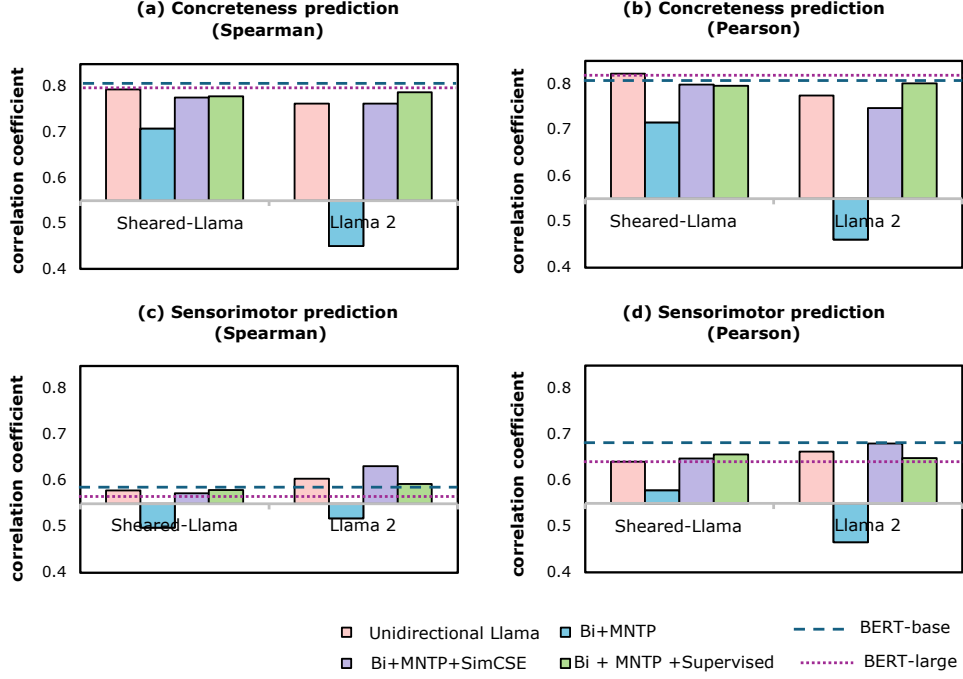


Figure 6: (a)-(b) show results for predicting noun concreteness (Task 3), while (c)-(d) display average results for predicting 11 sensorimotor dimensions of words (Task 4). Metrics include Spearman and Pearson correlation coefficients.

larger models.

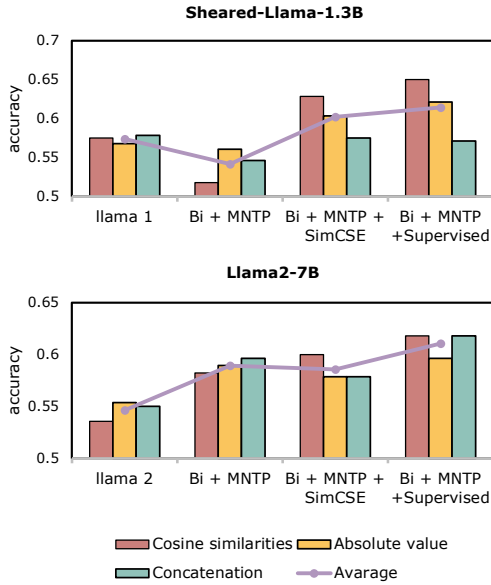


Figure 7: Results of word sense disambiguation task (Task 5). This task employs three methods of processing two target words' representation from pair sentences: cosine similarity, absolute difference, and concatenation. The metric is accuracy.

Finding 3: In the sense disambiguation task, contrastive learning methods improve the quality of embeddings from autoregressive models irrespec-

tive of the strategy for extracting probe features.

Given that **Word-in-Context** requires, compared to the other datasets, the combination of two contextualized representations, we tested three different strategies to combine the vectors (Figure 7). No big differences can be seen between those strategies with the base models, and in this case Bi + MNTP has a mixed effect, leading to the usual performance deterioration with Sheared-Llama and to better performance with Llama 2. Interestingly, contrastive learning leads again to best scores in the task, and it can be noticed that after its application the cosine similarities strategy becomes the most consistent one (cosine strategy with Sheared-Llama and either one of contrastive training types are the only combinations outperforming both the BERT baselines). As this strategy is simply based on using vector proximity as the only feature for the probe, the result suggests that refining representations for sentence-level tasks with contrastive learning also leads to more meaningful distances between token-level embeddings in the vector space.

4.1 Anisotropy in the Embedding Space

Anisotropy, a known issue in pre-trained language models, has been attributed to the disproportionate influence of rare tokens in the negative direction of

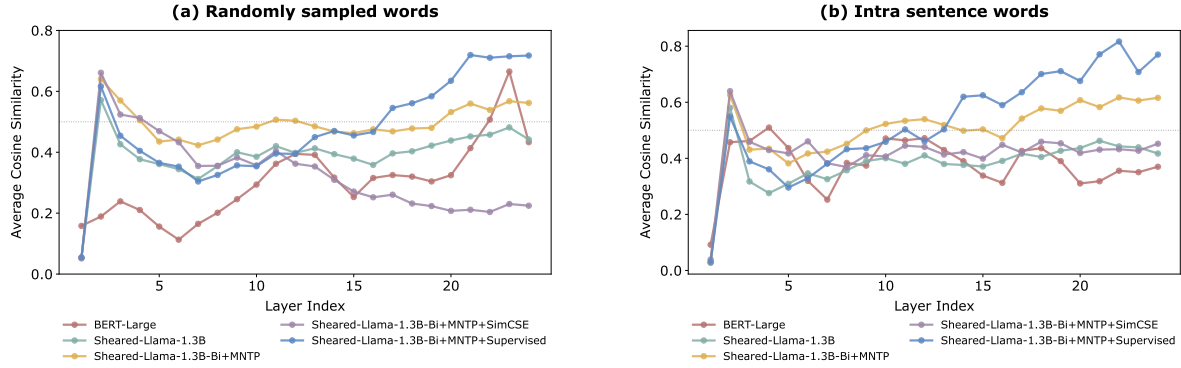


Figure 8: (a) represents the experimental results with randomly sampled words to examine the anisotropy; (b) illustrates the variation in anisotropy of words within the same sentence as the layer depth of the LLMs increases.

hidden states within likelihood-maximizing models (Wang et al., 2020a; Gao et al., 2019). However, in the literature, bidirectional attention models were shown to exhibit lower anisotropy compared to autoregressive models (Ethayarajh, 2019), possibly due to the impact of the attention mechanisms. Based on this, we speculate that bidirectional attention may alleviate the issue of anisotropy in the embeddings of autoregressive models.

Therefore, we extract the embeddings for the target words in the WIC dataset (Pilehvar and Camacho-Collados, 2019), randomly sample the representation of all target words 1000 times, and then calculate the average cosine similarity. The result is presented in Figure 8 a): neither BERT nor Llama produces isotropic word representations. Additionally, all Llama models have a sudden sharp increase in the cosine similarities in the second layer, followed by a decrease.

In terms of broad trends from shallow to deep layers, Llama models are relatively stable in the central layers and then the similarities increase again towards the last layers, with the exception for unsupervised contrastive learning Llama. The observed impact of the bidirectional attention mechanism contradicts our hypothesis: Sheared-Llama-1.3B with bidirectional attention exhibits higher global similarity across all hidden states compared to its unidirectional counterpart. While the model combining bidirectional attention and supervised contrastive learning is the one with the highest degree of anisotropy, the version with unsupervised contrastive learning is the most successful in reducing average similarities, showing an overall decreasing trend in the later layers. As for BERT, the similarities show a constant upward trend while moving from the earlier to the later layers, which

align with the previous analysis, and showing a spike in the degree of anisotropy in the very last ones (Ethayarajh, 2019).

On the basis of such findings, we would recommend that for tasks involving an unsupervised evaluation based on vector-space similarity, researchers adopt LLM with bidirectional training and unsupervised contrastive learning, since it seems to be the most robust combination against the anisotropy issue. On the other hand, our results also show that highly anisotropic representations do not necessarily have a negative impact on supervised tasks. In Figure 8 a) it can be clearly seen that, whereas unsupervised contrastive learning reduces anisotropy, supervised contrastive learning increases it. However, our experiments on the probing datasets showed that Llama models trained with supervised contrastive learning generally outperform all the other models.

To investigate how word representations within the same sentence evolve from shallow to deep layers of the model, we also extract words from individual sentences and perform layer-wise cosine similarity calculations to quantify intra-sentence anisotropy. Figure 8 b) shows that supervised contrastive learning bidirectional and bidirectional-only Llama models exhibit increasing anisotropy across layers, indicating that, in the deeper layers of the model, words from the same sentence tend to converge towards similar representations. This exhibits a trend of increasing anisotropy analogous to what we observed with randomly selected words. In the case of intra-sentence words, however, unsupervised contrastive learning does not significantly decrease the similarities compared to the base model, and the lowest levels of anisotropy are generally achieved by BERT in the late layers.

5 Conclusion

This work investigates how bidirectional attention contributes to contextual information encoding in word representations. The main contributions are: 1) providing a fine-grained word-based analysis of LLMs with enabled bidirectional attention, 2) unveiling the trade-off brought by bidirectional attention, and 3) contributing valuable insight into the application of LLM2Vec and the ongoing exploration of using LLMs as text encoders.

Contrary to our expectations, we found that in semantic probing tasks simply enabling bidirectional attention is not sufficient, as it may decrease the model’s capability to represent the previous context. However, we observed that across tasks, with the only exception of regression tasks for norms prediction, adding contrastive learning on top of bidirectional attention tends to improve the representation quality of the embeddings extracted from autoregressive LLMs. The fact that the sense disambiguation task of the *Word-in-Context* dataset, after contrastive learning, can be better addressed simply by using the cosine between the vectors suggests that this technique might lead to more meaningful distances between token-level vectors.

Moreover, we further analyzed the anisotropy of embedding representations across various hidden states. We found that bidirectional attention increases representation isotropy across all Llama layers, showing that this feature is not inherently related to unidirectional attention. Among the contrastive learning strategies, the unsupervised one seems to improve the anisotropy issue in the vector space, and thus it would be the recommended choice if the goal was to adapt an autoregressive LLM to unsupervised word embedding tasks that are evaluated via vector similarity.

Limitations

This probing study reveals that activating bidirectional attention impacts the semantic encoding of word representations, though the underlying mechanisms remain unclear. Additionally, the correspondence between high-dimensional dense vectors in computational models and semantic information warrants further investigation to bridge interpretability gaps in neural representation studies.

Another limitation is that our study is limited to the English language with the Sheared-Llama-1.3B and Llama2-7B models, and bigger models were not tested due to limitation of our computational

resources. Different languages and models may yield varying effects on the performance our lexical semantic tasks. We anticipate that future work can expand to diverse languages and models to validate and refine our findings. Likewise, Our approach to extracting contextual word representations, while designed for broad model compatibility, may not be necessarily the optimal one for each model.

Lastly, while bidirectional attention mechanisms have been shown to enhance performance in text embedding tasks in LLMs, their inherent capabilities were not thoroughly evaluated in our study. However, since our paper focuses on the problem of adapting autoregressive LLMs to embedding-based tasks, we did not evaluate text generation capabilities. To our knowledge, bidirectional attention and contrastive learning are not commonly adopted for text-generation tasks. A previous study (Khosla et al., 2025) demonstrated that enabling bidirectional attention in LLMs may significantly exacerbate text repetition. This occurs because such mechanisms disrupt the LLM’s native autoregressive generation, causing it to resemble BERT-like models, which are known to struggle with coherent text production. Future work may include more in-depth explorations of this problem.

Acknowledgements

We sincerely thank Behnam Ghader et al. (2024) for their innovative work on LLM2Vec, which provided foundational inspiration for our research. Their open-source contribution of the models on HuggingFace have been invaluable to our work. We also thank anonymous reviewers for their constructive feedback and suggestions, which have significantly helped us to refine this study.

Ethics Statement

This study investigates the impact of bidirectional attention mechanisms on semantic feature extraction in LLMs. While the technical contributions aim to advance model interpretability and linguistic capability, we acknowledge broader ethical responsibilities in disseminating and applying these findings. The research is intended to support beneficial applications in natural language understanding, such as improving machine translation and text summarization, where enhanced semantic analysis could yield significant societal value.

References

- Parishad BehnamGhader, Vaibhav Adlakha, Marius Mosbach, Dzmitry Bahdanau, Nicolas Chapados, and Siva Reddy. 2024. LLM2Vec: Large language models are secretly powerful text encoders. In *Proceedings of COLM*.
- Rishi Bommasani, Kelly Davis, and Claire Cardie. 2020. [Interpreting Pretrained Contextualized Representations via Reductions to Static Embeddings](#). In *Proceedings of ACL*, pages 4758–4781, Online. Association for Computational Linguistics.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33:1877–1901.
- John A Bullinaria and Joseph P Levy. 2012. Extracting semantic representations from word co-occurrence statistics: stop-lists, stemming, and SVD. *Behavior Research Methods*, 44:890–907.
- Xingyu Cai, Jiaji Huang, Yuchen Bian, and Kenneth Church. 2021. Isotropy in the contextual embedding space: Clusters and manifolds. In *International conference on learning representations*.
- Emmanuele Chersoni, Enrico Santus, Chu-Ren Huang, Alessandro Lenci, et al. 2021. Decoding word embeddings with brain-based semantic features. *Computational Linguistics*, 47(3):663–698.
- Emmanuele Chersoni, Rong Xiang, Qin Lu, and Chu-Ren Huang. 2020. Automatic Learning of Modality Exclusivity Norms with Crosslingual Word Embeddings. In *Proceedings of *SEM*.
- Gabriella Chronis and Katrin Erk. 2020. [When is a bishop not like a rook? when it’s like a rabbi! multi-prototype BERT embeddings for estimating semantic relationships](#). In *Proceedings of CONLL*, pages 227–244, Online. Association for Computational Linguistics.
- Kevin Clark, Minh-Thang Luong, Quoc V. Le, and Christopher D. Manning. 2020. [ELECTRA: Pre-training text encoders as discriminators rather than generators](#). In *International Conference on Learning Representations*.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. [BERT: Pre-training of deep bidirectional transformers for language understanding](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4171–4186, Minneapolis, Minnesota. Association for Computational Linguistics.
- Kawin Ethayarajh. 2019. [How contextual are contextualized word representations? Comparing the geometry of BERT, ELMo, and GPT-2 embeddings](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 55–65, Hong Kong, China. Association for Computational Linguistics.
- Luana Fagarasan, Eva Maria Vecchi, and Stephen Clark. 2015. From Distributional Semantics to Feature Norms: Grounding Semantic Models in Human Perceptual Data. In *Proceedings of IWCS*.
- Manaal Faruqui, Yulia Tsvetkov, Pushpendre Rastogi, and Chris Dyer. 2016. [Problems with evaluation of word embeddings using word similarity tasks](#). In *Proceedings of the 1st Workshop on Evaluating Vector-Space Representations for NLP*, pages 30–35, Berlin, Germany. Association for Computational Linguistics.
- Michael M Flor. 2024. Three Studies on Predicting Word Concreteness with Embedding Vectors. In *Proceedings of the LREC-COLING Workshop on Cognitive Aspects of the Lexicon*.
- Jun Gao, Di He, Xu Tan, Tao Qin, Liwei Wang, and Tieyan Liu. 2019. [Representation degeneration problem in training natural language generation models](#). In *International Conference on Learning Representations*.
- Tianyu Gao, Xingcheng Yao, and Danqi Chen. 2021. [SimCSE: Simple contrastive learning of sentence embeddings](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6894–6910, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Anna Gladkova and Aleksandr Drozd. 2016. [Intrinsic evaluations of word embeddings: What can we do better?](#) In *Proceedings of the 1st Workshop on Evaluating Vector-Space Representations for NLP*, pages 36–42, Berlin, Germany. Association for Computational Linguistics.
- Nathan Godey, Éric de la Clergerie, and Benoît Sagot. 2024. Anisotropy is inherent to self-attention in transformers. In *Proceedings of EACL*.
- Lorenzo Gregori, Maria Montefinese, Daniele P Radicioni, Amelio Ravelli Andrea, Rossella Varvara, et al. 2020. CONCRETEXT@ EVALITA2020: The concreteness in context task. In *Ceur Workshop Proceedings*, volume 2765, pages 1–8. CEUR.
- Zellig S Harris. 1954. Distributional structure. *WORD*, 10(2-3):146–162.
- John Hewitt and Percy Liang. 2019. [Designing and interpreting probes with control tasks](#). In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 2733–2743, Hong Kong, China. Association for Computational Linguistics.

- John Hewitt and Christopher D. Manning. 2019. [A structural probe for finding syntax in word representations](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 4129–4138, Minneapolis, Minnesota. Association for Computational Linguistics.
- Carina Kauf, Anna A Ivanova, Giulia Rambelli, Emanuele Chersoni, Jingyuan Selena She, Zawad Chowdhury, Evelina Fedorenko, and Alessandro Lenci. 2023. Event Knowledge in Large Language Models: The Gap Between the Impossible and the Unlikely. *Cognitive Science*, 47(11):e13386.
- Savya Khosla, Aditi Tiwari, Kushal Kafle, Simon Jenni, Handong Zhao, John Collomosse, and Jing Shi. 2025. [Magnet: Augmenting generative decoders with representation learning and infilling capabilities](#). *Preprint*, arXiv:2501.08648.
- Josef Klafka and Allyson Ettinger. 2020. [Spying on your neighbors: Fine-grained probing of contextual embeddings for information about surrounding words](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4801–4811, Online. Association for Computational Linguistics.
- Chankyu Lee, Rajarshi Roy, Mengyao Xu, Jonathan Raiman, Mohammad Shoeybi, Bryan Catanzaro, and Wei Ping. 2025. [Nv-embed: Improved techniques for training llms as generalist embedding models](#). In *The Thirteenth International Conference on Learning Representations*.
- Dandan Li and Douglas Summers-Stay. 2019. Mapping Distributional Semantics to Property Morphs with Deep Neural Networks. *Big Data and Cognitive Computing*, 3(2):30.
- Tal Linzen, Emmanuel Dupoux, and Yoav Goldberg. 2016. [Assessing the ability of LSTMs to learn syntax-sensitive dependencies](#). *Transactions of the Association for Computational Linguistics*, 4:521–535.
- Nelson F. Liu, Matt Gardner, Yonatan Belinkov, Matthew E. Peters, and Noah A. Smith. 2019. [Linguistic knowledge and transferability of contextual representations](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 1073–1094, Minneapolis, Minnesota. Association for Computational Linguistics.
- Zhu Liu, Cunliang Kong, Ying Liu, and Maosong Sun. 2024. [Fantastic semantics and where to find them: Investigating which layers of generative LLMs reflect lexical semantics](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 14551–14558, Bangkok, Thailand. Association for Computational Linguistics.
- Gonzalo Martínez, Juan Diego Molero, Sandra González, Javier Conde, Marc Brysbaert, and Pedro Reviriego. 2025. Using large language models to estimate features of multi-word expressions: Concreteness, valence, arousal. *Behavior Research Methods*, 57(1):1–11.
- Jacob Matthews, John Starr, and Marten Schijndel. 2024. [Semantics or spelling? probing contextual word embeddings with orthographic noise](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 4495–4504, Bangkok, Thailand. Association for Computational Linguistics.
- Eleni Metheniti, Tim Van De Cruys, and Nabil Hathout. 2022. [About time: Do transformers learn temporal verbal aspect?](#) In *Proceedings of the Workshop on Cognitive Modeling and Computational Linguistics*, pages 88–101, Dublin, Ireland. Association for Computational Linguistics.
- Tomas Mikolov, Kai Chen, Greg Corrado, and Jeffrey Dean. 2013. [Efficient estimation of word representations in vector space](#). *Preprint*, arXiv:1301.3781.
- Niklas Muennighoff. 2022. SGPT: GPT sentence embeddings for semantic search. *arXiv preprint arXiv:2202.08904*.
- Humza Naveed, Asad Ullah Khan, Shi Qiu, Muhammad Saqib, Saeed Anwar, Muhammad Usman, Naveed Akhtar, Nick Barnes, and Ajmal Mian. 2023. A comprehensive overview of large language models. *arXiv preprint arXiv:2307.06435*.
- Arvind Neelakantan, Tao Xu, Raul Puri, Alec Radford, Jesse Michael Han, Jerry Tworek, Qiming Yuan, Nikolas Tezak, Jong Wook Kim, Chris Hallacy, et al. 2022. Text and code embeddings by contrastive pre-training. *arXiv preprint arXiv:2201.10005*.
- Jeffrey Pennington, Richard Socher, and Christopher Manning. 2014. [GloVe: Global vectors for word representation](#). In *Proceedings of the 2014 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543, Doha, Qatar. Association for Computational Linguistics.
- Matthew E. Peters, Mark Neumann, Mohit Iyyer, Matt Gardner, Christopher Clark, Kenton Lee, and Luke Zettlemoyer. 2018. [Deep contextualized word representations](#). In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 2227–2237, New Orleans, Louisiana. Association for Computational Linguistics.
- Mohammad Taher Pilehvar and Jose Camacho-Collados. 2019. [WiC: the word-in-context dataset for evaluating context-sensitive meaning representations](#). In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 1267–1273, Minneapolis, Minnesota. Association for Computational Linguistics.

- Muhammad Qorib, Geonsik Moon, and Hwee Tou Ng. 2024. [Are decoder-only language models better than encoder-only language models in understanding word meaning?](#) In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 16339–16347, Bangkok, Thailand. Association for Computational Linguistics.
- Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. 2019. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9.
- Hinrich Schütze. 1992. [Dimensions of Meaning](#). In *Supercomputing '92: Proceedings of the 1992 ACM/IEEE Conference on Supercomputing*, pages 787–796.
- Jacob Mitchell Springer, Suhas Kotha, Daniel Fried, Graham Neubig, and Aditi Raghunathan. 2025. [Repetition improves language model embeddings](#). In *The Thirteenth International Conference on Learning Representations*.
- Bill Thompson and Gary Lupyan. 2018. Automatic Estimation of Lexical Concreteness in 77 Languages. In *Proceedings of the Annual Meeting of the Cognitive Science Society*, volume 40.
- William Timkey and Marten van Schijndel. 2021. [All bark and no bite: Rogue dimensions in transformer language models obscure representational quality](#). In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 4527–4546, Online and Punta Cana, Dominican Republic. Association for Computational Linguistics.
- Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. 2023a. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971*.
- Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajwal Bhargava, Shruti Bhosale, et al. 2023b. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*.
- Sean Trott and Benjamin Bergen. 2022. [Contextualized sensorimotor norms: multi-dimensional measures of sensorimotor strength for ambiguous english words, in context](#). *Preprint*, arXiv:2203.05648.
- Akira Utsumi. 2018. A Neurobiologically Motivated Analysis of Distributional Semantic Models. In *Proceedings of CogSci*.
- Elena Voita and Ivan Titov. 2020. [Information-theoretic probing with minimum description length](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 183–196, Online. Association for Computational Linguistics.
- Hetong Wang, Pasquale Minervini, and Edoardo Ponti. 2024a. [Probing the emergence of cross-lingual alignment during LLM training](#). In *Findings of the Association for Computational Linguistics: ACL 2024*, pages 12159–12173, Bangkok, Thailand. Association for Computational Linguistics.
- Liang Wang, Nan Yang, Xiaolong Huang, Linjun Yang, Rangan Majumder, and Furu Wei. 2024b. [Improving text embeddings with large language models](#). In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 11897–11916, Bangkok, Thailand. Association for Computational Linguistics.
- Lingxiao Wang, Jing Huang, Kevin Huang, Ziniu Hu, Guangtao Wang, and Quanquan Gu. 2020a. [Improving neural language generation with spectrum control](#). In *International Conference on Learning Representations*.
- Yuxuan Wang, Yutai Hou, Wanxiang Che, and Ting Liu. 2020b. From static to dynamic word representations: a survey. *International Journal of Machine Learning and Cybernetics*, 11:1611–1630.
- Zhiyong Wu, Yun Chen, Ben Kao, and Qun Liu. 2020. [Perturbed masking: Parameter-free probing for analyzing and interpreting BERT](#). In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4166–4176, Online. Association for Computational Linguistics.
- Mengzhou Xia, Tianyu Gao, Zhiyuan Zeng, and Danqi Chen. 2023. Sheared Llama: Accelerating language model pre-training via structured pruning. *arXiv preprint arXiv:2310.06694*.
- Qihui Xu, Yingying Peng, Samuel A Nastase, Martin Chodorow, Minghua Wu, and Ping Li. 2023. Does conceptual representation require embodiment? insights from large language models. *arXiv preprint arXiv:2305.19103*.
- Yilun Zhu, Joel Ruben Antony Moniz, Shruti Bhargava, Jiarui Lu, Dhivya Piraviperumal, Site Li, Yuan Zhang, Hong Yu, and Bo-Hsiang Tseng. 2024. [Can large language models understand context?](#) In *Findings of the Association for Computational Linguistics: EACL 2024*, pages 2004–2018, St. Julian’s, Malta. Association for Computational Linguistics.

A Experimental Setup

A.1 Dataset Details

Table 1 details datasets used in experiments. The hyperlinks are attached to dataset names, which direct to the download page of each dataset.

A.2 Model Details

Table 3 shows used model details in the experiments. Hyperlinks to each model’s Hugging Face page are added.

A.3 Other Setup Details

The hyperparameters of the Multilayer Perceptron (MLP) classifier used in our experiments are presented in Table 2. All experiments were implemented in Python 3.10 utilizing PyTorch 2.6.0 (CUDA 12.4), Transformers 4.43.1, and PEFT 0.10.0 libraries. To ensure reproducibility across experiments, we set the random seed to a fixed value of 42. To reduce memory usage and computational cost, we employed half-precision (FP16) arithmetic for embedding extraction.

B Experimental Results

B.1 Multilayer Perceptron

Tables 4 to 7 shows the experimental results of five tasks using multilayer perceptron as the probe model.

B.2 Logistic Regression/Linear Regression

Tables 8 to 11 shows the experimental results of five tasks using logistic regression or linear regression as the probe model.

C Effectiveness of Probing Method

As Voita and Titov (2020) pointed out, the probe accuracy can be similar when probing for genuine linguistic labels and probing for random synthetic tasks. To address this concern, we introduced control tasks to evaluate the selectivity of the probes. In control tasks, each label is not genuine but randomly sampled. As Hewitt and Liang (2019) suggested, a good probe should achieve higher accuracy on linguistic tasks and lower accuracy on control tasks. We test BERT-base, Sheared-Llama, and Sheared-Llama (Bi + BNTP) on control tasks (MLP classifier).

Results are shown in Tables 12 to 14: Across all five tasks we examined, the control task consistently showed significantly lower accuracy than the original task, demonstrating that the probes exhibited high selectivity in all five cases.

D Task Sensitivity to Context

To assess the importance of the context in the first four tasks (i.e., whether performance benefits from contextual embeddings), we conducted additional experiments using non-contextual inputs (probed words without surrounding context) and MLP classifier. As demonstrated in Tables 15 and 16, the

results confirm that all four tasks indeed show improved performance with contextual information.

In Task 1 (Table 15): for each experiment, we input only a single word (the probed word) to the model, rather than an entire sentence. As expected, we found that accuracy is high only when probing the target word directly, with the other words yielding random-guess performance. Notice that in Task 1 our focus is not on target word probing accuracy itself, but rather on how much semantic information about the target word is contained in the embeddings of the other words. Therefore, task 1 can be solved only by contextualized embeddings for an input token other than the probed word.

As for Task 2, Task 3, and Task 4, we also added no-context results (Table 16). We found that results without context are significantly lower than those with context, indicating that contextual information is crucial for prediction in all three tasks.

Dataset	Predictive Method	Train Data	Test Data
Context-probe	Classification	Subject: 4000 sentences Object: 4000 sentences Verb: 8000 sentences	Subject: 1000 sentences Object: 1000 sentences Verb: 2000 sentences
Lexical aspect binary Classification (telicity/duration)	Classification	Telicity: 3920 sentences Duration: 2591 sentences	Telicity: 980 sentences Duration: 648 sentences
CONcreTEXT	Regression	347 sentences	87 sentences
Contextualized Sensorimotor Norms	Regression	358 sentences	90 sentences
WiC Dataset	Classification	1120 sentence pairs	280 sentence pairs

Table 1: Dataset details and train-test data separation of classifier or regressor training in our experiments. The predictive method applied and data separation between the train and test dataset are demonstrated. Links leading to the webpage of Github or the related recourse are added to the dataset name.

	T1	T2-T	T2-D	T3	T4	T5-A	T5-S	T5-C
Structure	[20, 10]	[20, 10]	[20, 10]	[40, 20]	[40, 20]	[20,10]	[5,2]	[20, 10]
Batch size	16	16	16	8	16	16	8	16
Learning rate	2e-5	2e-5	2e-5	2e-3	2e-3	1e-5	1e-3	1.5e-5

Table 2: Hyperparameters in Multilayer Perceptron classifier in experiments. All language models utilize the same hyperparameters setting in the same task. The first row indicates Task 1, Task 2 telicity, Task 3 duration, Task 3, Task 4, Task 5 (absolute difference), Task 5 (cosine similarity), and Task 5 (concatenation).

Family	Model Link	Parameters	Attention Mechanism
BERT	bert-base-uncased	109M	Bidirectional
	bert-large-uncased	335M	
Llama 1	Sheared-LLaMA-1.3B	1.3B	Unidirectional
	LLM2Vec-Sheared-LLaMA-mntp		Bidirectional
	LLM2Vec-Sheared-LLaMA-mntp-unsup-simcse		
	LLM2Vec-Sheared-LLaMA-mntp-supervised		
Llama 2	Llama-2-7b-hf	7B	Unidirectional
	LLM2Vec-Llama-2-7b-chat-hf-mntp		Bidirectional
	LLM2Vec-Llama-2-7b-chat-hf-mntp-unsup-simcse		
	LLM2Vec-Llama-2-7b-chat-hf-mntp-supervised		

Table 3: Model details of used models, parameters details, and enabled attention mechanism in our experiments. Links leading to the Hugging Face page of each model are added to the hyperlinks.

Model	Subtask	Index 1	Index 2	Index 3	Index 4	Index 5
BERT-base	subject animacy	0.942	1.000	0.966	0.944	0.944
	verb causative	0.989	0.976	0.999	0.990	0.989
	verb dynamic	0.998	0.979	1.000	0.999	0.984
	object animacy	0.942	0.870	0.914	0.948	0.982
BERT-large	subject animacy	0.967	0.998	0.965	0.969	0.939
	verb causative	0.989	0.945	0.998	0.990	0.968
	verb dynamic	0.992	0.975	1.000	0.991	0.971
	object animacy	0.932	0.847	0.889	0.940	0.987
Sheared-Llama-1.3B	subject animacy	0.547†	1.000	0.994	0.996	0.943
	verb causative	0.611†	0.534†	1.000	0.995	0.987
	verb dynamic	0.515†	0.494†	1.000	0.998	0.973
	object animacy	0.539†	0.501†	0.508†	0.507†	1.000
Sheared-Llama-1.3B (Bi + MNTP)	subject animacy	0.976	0.976★	0.916★	0.826★	0.766★
	verb causative	0.958	0.968	0.991★	0.960★	0.939★
	verb dynamic	0.972	0.974	0.983★	0.954★	0.935★
	object animacy	0.850	0.877	0.960	0.995	0.999★
Sheared-Llama-1.3B (Bi + MNTP + SimCSE)	subject animacy	0.995	1.000	0.949	0.949	0.957
	verb causative	0.993	0.999	1.000	0.995	0.997
	verb dynamic	0.999	1.000	1.000	0.997	0.997
	object animacy	0.944	0.956	0.978	0.996	1.000
Sheared-Llama-1.3B (Bi + MNTP + Supervised)	subject animacy	0.988	1.000	0.980	0.983	0.993
	verb causative	0.998	0.995	1.000	0.999	0.999
	verb dynamic	1.000	0.999	1.000	0.999	0.997
	object animacy	0.993	0.995	0.989	0.999	1.000
Llama2-7B	subject animacy	0.547†	1.000	0.995	0.996	0.962
	verb causative	0.611†	0.526†	1.000	1.000	0.997
	verb dynamic	0.515†	0.515†	1.000	1.000	0.997
	object animacy	0.544†	0.496†	0.497†	0.510†	1.000
Llama2-7B (Bi + MNTP)	subject animacy	0.999	0.998★	0.978★	0.908★	0.897★
	verb causative	0.979	0.990	1.000	0.986★	0.978★
	verb dynamic	0.985	0.990	0.996★	0.983★	0.978★
	object animacy	0.901	0.941	0.994	0.999	0.997★
Llama2-7B (Bi + MNTP + SimCSE)	subject animacy	0.993	1.000	0.984	0.982	0.994
	verb causative	0.998	1.000	1.000	0.999	0.999
	verb dynamic	0.998	1.000	1.000	0.999	1.000
	object animacy	0.991	0.996	0.988	0.996	1.000
Llama2-7B (Bi + MNTP + Supervised)	subject animacy	1.000	1.000	0.991	0.987	0.999
	verb causative	1.000	0.999	1.000	0.999	0.999
	verb dynamic	1.000	1.000	1.000	1.000	0.997
	object animacy	0.983	0.990	0.986	0.999	1.000

Table 4: Results of Task 1 (context probes, Multilayer Perceptron). Numbers with † give evidence that LLMs can not access the subsequent context, and numbers with ★ show that bidirectional attention weakens preceding context utilization.

Model	T2-T	T2-D	T5-A	T5-S	T5-C	T5-Avg
BERT-base	0.8286	0.9167	0.6071	0.6179	0.5857	0.6036
BERT-large	0.8235	0.9059	0.6429	0.6786	<u>0.5393</u>	0.6203
Sheared-Llama-1.3B	0.8235	0.9074	0.5679	0.5750	0.5786	0.5738
Sheared-Llama-1.3B-Bi+MNTP	0.7653	0.8750	0.5607	<u>0.5179</u>	0.5464	<u>0.5417</u>
Sheared-Llama-1.3B-Bi+MNTP+SimCSE	0.8276	0.9120	0.6036	0.6286	0.5750	0.6024
Sheared-Llama-1.3B-Bi+MNTP+Supv.	0.8265	0.9198	0.6214	0.6500	0.5714	0.6143
Llama2-7B	0.8388	0.9090	<u>0.5536</u>	0.5357	0.5500	0.5464
Llama2-7B-Bi+MNTP	<u>0.7316</u>	<u>0.8596</u>	0.5893	0.5821	0.5964	0.5893
Llama2-7B-Bi+MNTP+SimCSE	0.8531	0.9306	0.5786	0.6000	0.5786	0.5857
Llama2-7B-Bi+MNTP+Supv.	0.8531	0.9352	0.5964	0.6179	0.6179	0.6107

Table 5: Classification accuracy of Task 2 and Task 5 (Multilayer Perceptron): T2-T (Task 2 telicity subtask), T2-D (Task 2 duration subtask), T5-A/S/C (Task 5 with the absolute difference/cosine similarity/concatenation.) “Supv.” is the abbreviation of “Supervised”. The bolded numbers indicate the highest performances among models in the same column, while numbers with underlines mean the weakest performance.

Model	T3-MSE	T3-R ²	T3-Pearson	T3-Spearman
BERT-base	0.6758	0.6612	0.8215	0.7971
BERT-large	0.6639	0.6671	0.8219	0.7946
Sheared-Llama-1.3B	0.7336	0.6322	0.8219	0.7946
Sheared-Llama-1.3B-Bi+MNTP	0.9808	0.5082	0.7157	0.7079
Sheared-Llama-1.3B-Bi+MNTP+SimCSE	0.7383	0.6298	0.7983	0.7768
Sheared-Llama-1.3B-Bi+MNTP+Supv.	0.6926	0.6527	0.7958	0.7795
Llama2-7B	0.8213	0.5882	0.7751	0.7635
Llama2-7B-Bi+MNTP	<u>1.5857</u>	<u>0.2049</u>	<u>0.4616</u>	<u>0.4503</u>
Llama2-7B-Bi+MNTP+SimCSE	0.8872	0.5552	0.7477	0.7628
Llama2-7B-Bi+MNTP+Supv.	0.7241	0.6370	0.8012	0.7879

Table 6: Regression results of Task 3 (concreteness prediction, Multilayer Perceptron). “Supv.” is the abbreviation of “Supervised”.

Model	T4-MSE	T4-R ²	T4-Pearson	T4-Spearman
BERT-base	0.3738	0.4065	0.6706	0.5916
BERT-large	0.3905	0.3483	0.6300	0.5567
Sheared-Llama-1.3B	0.3884	0.3664	0.6408	0.5785
Sheared-Llama-1.3B-Bi+MNTP	0.4626	<u>-0.4610</u>	0.5788	<u>0.4976</u>
Sheared-Llama-1.3B-Bi+MNTP+SimCSE	0.3766	0.4215	0.6484	0.5730
Sheared-Llama-1.3B-Bi+MNTP+Supv.	0.3835	0.3798	0.6572	0.5798
Llama2-7B	0.3854	0.4154	0.6626	0.6040
Llama2-7B-Bi+MNTP	<u>0.5505</u>	0.1782	<u>0.4659</u>	0.5179
Llama2-7B-Bi+MNTP+SimCSE	0.3616	0.4240	0.6806	0.6317
Llama2-7B-Bi+MNTP+Supv.	0.3903	0.3841	0.6486	0.5927

Table 7: Average regression results of Task 4 (sensorimotor prediction, Multilayer Perceptron) “Supv.” is the abbreviation of “Supervised”.

Model	Subtask	Index 1	Index 2	Index 3	Index 4	Index 5
BERT-base	subject animacy	0.925	0.999	0.958	0.925	0.936
	verb causative	0.975	0.959	0.994	0.975	0.980
	verb dynamic	0.989	0.984	1.000	0.989	0.982
	object animacy	0.931	0.835	0.913	0.931	0.984
BERT-large	subject animacy	0.944	0.997	0.961	0.944	0.929
	verb causative	0.988	0.931	0.998	0.988	0.970
	verb dynamic	0.988	0.974	0.998	0.988	0.969
	object animacy	0.925	0.810	0.871	0.925	0.988
Sheared-Llama-1.3B	subject animacy	0.547 †	1.000	0.999	0.998	0.959
	verb causative	0.611 †	0.527 †	1.000	0.997	0.996
	verb dynamic	0.515 †	0.500 †	1.000	1.000	0.991
	object animacy	0.539 †	0.492 †	0.534 †	0.479 †	1.000
Sheared-Llama-1.3B (Bi + MNTP)	subject animacy	0.979	0.983 *	0.912 *	0.823 *	0.765 *
	verb causative	0.962	0.969	0.994 *	0.966 *	0.943 *
	verb dynamic	0.971	0.972	0.987 *	0.972 *	0.945 *
	object animacy	0.882	0.887	0.956	0.988	0.995 *
Sheared-Llama-1.3B (Bi + MNTP + SimCSE)	subject animacy	0.995	1.000	0.961	0.947	0.962
	verb causative	0.996	0.999	1.000	0.999	0.998
	verb dynamic	1.000	1.000	1.000	0.998	0.999
	object animacy	0.952	0.970	0.986	0.996	1.000
Sheared-Llama-1.3B (Bi + MNTP + Supervised)	subject animacy	0.970	1.000	0.990	0.991	0.994
	verb causative	0.986	0.998	1.000	1.000	0.998
	verb dynamic	0.996	1.000	1.000	1.000	0.999
	object animacy	0.994	0.997	0.992	0.999	1.000
Llama2-7B	subject animacy	0.547 †	1.000	0.998	0.998	0.977
	verb causative	0.611 †	0.528 †	1.000	1.000	0.994
	verb dynamic	0.512 †	0.490 †	1.000	1.000	0.996
	object animacy	0.544 †	0.496 †	0.507 †	0.505 †	0.999
Llama2-7B (Bi + MNTP)	subject animacy	0.997	0.993*	0.986*	0.949*	0.936*
	verb causative	0.979	0.990	0.997*	0.994*	0.983*
	verb dynamic	0.988	0.992	0.997*	0.990*	0.985*
	object animacy	0.932	0.958	0.989	0.993	0.994*
Llama2-7B (Bi + MNTP + SimCSE)	subject animacy	0.998	1.000	0.993	0.984	0.993
	verb causative	1.000	1.000	1.000	1.000	1.000
	verb dynamic	1.000	1.000	1.000	1.000	1.000
	object animacy	0.989	0.992	0.990	0.997	1.000
Llama2-7B (Bi + MNTP + Supervised)	subject animacy	1.000	1.000	0.993	0.988	0.999
	verb causative	1.000	0.999	1.000	1.000	0.998
	verb dynamic	1.000	1.000	1.000	0.999	0.999
	object animacy	0.982	0.991	0.989	0.999	1.000

Table 8: Results of Task 1 (context probes, Logistic Regression). Numbers with † give evidence that LLMs can not access the subsequent context, and numbers with * show that bidirectional attention weakens preceding context utilization.

Model	T2-T	T2-D	T5-A	T5-S	T5-C	T5-Avg
BERT-base	0.7765	0.8904	0.5929	0.6179	0.5607	0.5905
BERT-large	0.7837	0.8796	0.5357	0.6643	<u>0.5000</u>	0.5666
Sheared-Llama-1.3B	0.7571	0.9059	0.4964	0.5357	0.5107	0.5142
Sheared-Llama-1.3B-Bi+MNTP	0.7010	0.8611	0.5321	0.5500	0.5071	0.5297
Sheared-Llama-1.3B-Bi+MNTP+SimCSE	0.7745	0.9090	0.5500	0.6321	0.5357	0.5726
Sheared-Llama-1.3B-Bi+MNTP+Supv.	0.7918	0.9136	0.5821	0.6429	0.5500	0.5917
Llama2-7B	0.8061	0.9151	<u>0.4679</u>	<u>0.5107</u>	<u>0.5000</u>	<u>0.4929</u>
Llama2-7B-Bi+MNTP	<u>0.6704</u>	<u>0.8596</u>	0.5179	0.5500	0.5250	0.5310
Llama2-7B-Bi+MNTP+SimCSE	0.8194	0.9244	0.5571	0.6000	0.5429	0.5667
Llama2-7B-Bi+MNTP+Supv.	0.8357	0.9198	0.5500	0.6393	0.5321	0.5738

Table 9: Classification accuracy of Task 2 and Task 5 (Logistic Regression): T2-T (Task 2 telicity subtask), T2-D (Task 2 duration subtask), T5-A/S/C (Task 5 with the absolute difference/cosine similarity/concatenation.) “Supv.” is the abbreviation of “Supervised”.

Model	T3-MSE	T3-R ²	T3-Pearson	T3-Spearman
BERT-base	1.2702	0.3631	0.6967	0.6774
BERT-large	0.9944	0.5014	0.7457	0.7319
Sheared-Llama-1.3B	0.8761	0.5607	0.7778	0.7588
Sheared-Llama-1.3B-Bi+MNTP	1.1636	0.4166	0.6866	0.6792
Sheared-Llama-1.3B-Bi+MNTP+SimCSE	0.7638	0.6171	0.7914	0.7668
Sheared-Llama-1.3B-Bi+MNTP+Supv.	0.7269	0.6356	0.8107	0.7905
Llama2-7B	0.8379	0.5799	0.7804	0.7638
Llama2-7B-Bi+MNTP	<u>1.5123</u>	<u>0.2418</u>	<u>0.5370</u>	<u>0.5435</u>
Llama2-7B-Bi+MNTP+SimCSE	0.8108	0.5935	0.7895	0.7724
Llama2-7B-Bi+MNTP+Supv.	0.6597	0.6692	0.8375	0.8190

Table 10: Regression results of Task 3 (concreteness prediction, Linear Regression). “Supv.” is the abbreviation of “Supervised”.

Model	T4-MSE	T4-R ²	T4-Pearson	T4-Spearman
BERT-base	<u>0.6625</u>	0.0215	0.5680	0.4917
BERT-large	0.5907	0.0862	0.5695	0.5100
Sheared-Llama-1.3B	0.4152	0.3446	0.6484	0.5859
Sheared-Llama-1.3B-Bi+MNTP	0.6371	0.0201	0.5185	0.4603
Sheared-Llama-1.3B-Bi+MNTP+SimCSE	0.4029	0.3829	0.6662	0.6085
Sheared-Llama-1.3B-Bi+MNTP+Supv.	0.4554	0.2950	0.6304	0.5622
Llama2-7B	0.3827	0.3994	0.6636	0.5944
Llama2-7B-Bi+MNTP	0.6598	<u>0.0121</u>	<u>0.4925</u>	<u>0.4301</u>
Llama2-7B-Bi+MNTP+SimCSE	0.3560	0.4508	0.6952	0.6327
Llama2-7B-Bi+MNTP+Supv.	0.3547	0.4577	0.6965	0.6271

Table 11: Average regression results of Task 4 (sensorimotor prediction, Linear Regression). “Supv.” is the abbreviation of “Supervised”.

Model	Task	Index1	Index2	Index3	Index4	Index5
BERT-base	control	0.518	0.494	0.529	0.519	0.520
	original	0.942	1.000	0.966	0.944	0.944
Sheared-Llama	control	0.497	0.520	0.490	0.530	0.510
	original	0.547	1.000	0.994	0.996	0.943
Sheared-Llama (Bi+BNTP)	control	0.510	0.527	0.485	0.503	0.493
	original	0.976	0.976	0.916	0.826	0.766

Table 12: This table presents a comparison of accuracy between the control task and the original task in Task 1 (subject animacy).

Model	Task	Task3-P	Task3-S	Task4-P	Task4-S
BERT-base	control	0.0649	0.0810	0.0421	0.0417
	original	0.8215	0.7971	0.6706	0.5916
Sheared-Llama	control	-0.0798	-0.0739	-0.0097	-0.0091
	original	0.8219	0.7946	0.6408	0.5785
Sheared-Llama (Bi + BNTP)	control	0.1214	0.1228	0.0437	0.0337
	original	0.7157	0.7079	0.5788	0.4976

Table 13: This table presents a comparison of Pearson and Spearman correlation coefficient between the control task and the original task in Task 3 and 4. P and S represent Pearson and Spearman, respectively.

Model	Task	T2-T	T2-D	T5-A	T5-S	T5-C	T5-Avg
BERT-base	control	0.5173	0.5015	0.5357	0.5036	0.5179	0.5191
	original	0.8286	0.9167	0.6071	0.6179	0.5857	0.6036
Sheared-llama	control	0.5133	0.5448	0.5607	0.4750	0.5321	0.5226
	original	0.8235	0.9074	0.5679	0.5750	0.5786	0.5738
Sheared-llama (Bi + BNTP)	control	0.5163	0.5509	0.5143	0.5000	0.5179	0.5107
	original	0.7653	0.8750	0.5607	0.5179	0.5464	0.5417

Table 14: This table presents a comparison of accuracy between the control task and the original task in Task 2 and 5. T2-T (Task 2 telicity subtask), T2-D (Task 2 duration subtask), T5-A/S/C (Task 5 with the absolute difference/cosine similarity/concatenation).

Model	Subtask	Index1	Index2	Index3	Index4	Index5
BERT-base	subject animacy	0.507	1.000	0.519	0.507	0.527
	verb causative	0.511	0.527	1.000	0.508	0.510
	verb dynamic	0.516	0.511	1.000	0.516	0.530
	object animacy	0.515	0.533	0.501	0.482	1.000
Sheared-Llama	subject animacy	0.482	1.000	0.491	0.482	0.499
	verb causative	0.511	0.532	1.000	0.511	0.518
	verb dynamic	0.515	0.494	1.000	0.515	0.509
	object animacy	0.507	0.498	0.501	0.507	1.000
Sheared-Llama (Bi + MNTP)	subject animacy	0.518	1.000	0.493	0.518	0.494
	verb causative	0.511	0.535	1.000	0.511	0.522
	verb dynamic	0.515	0.497	1.000	0.485	0.513
	object animacy	0.507	0.514	0.498	0.507	1.000

Table 15: This table presents the accuracy of three semantic features in Task 1 without giving context.

Model	Context	Task2-T	Task2-D	Task3- P	T3- S	Task4- P	T4- S
BERT-base	with	0.8286	0.9167	0.8215	0.7971	0.6706	0.5916
	without	0.7480	0.8302	0.7105	0.6874	0.5405	0.5015
Sheared-Llama	with	0.8235	0.9074	0.8219	0.7946	0.6408	0.5785
	without	0.7531	0.8364	0.7270	0.7075	0.5880	0.5403
Sheared-Llama (Bi + MNTP)	with	0.7653	0.8750	0.7157	0.7079	0.5788	0.4976
	without	0.7367	0.8318	0.7034	0.6830	0.5714	0.5278

Table 16: This table presents a comparison of accuracy in Task 2, 3, 4 and 5 with and without giving context. T and D in Task 2 denote Telicity and Duration, and P and S in Task 3 and 4 represent the Pearson and Spearman.